

# Spark Walmart STOCK Data Analysis Project

## 1. Start a simple Spark Session

```
In [1]: import findspark
findspark.init()
import pyspark

from pyspark.sql import SparkSession

In [3]: spark=SparkSession.builder.appName('Big data Hadoop Spark Mini-Project').getOrCreate()

In [4]: sc = spark.sparkContext

In [5]: sc

Out[5]: SparkContext

Spark UI
Version
v3.3.2
Master
local[*]
AppName
Python Spark SQL Lesson
```

## 2. Load the Walmart Stock CSV File, have Spark infer the data types. What are the column names?

```
: walmartDF = spark.read.load('walmart_stock.csv', format='csv', sep=',',
inferSchema='true', header='true')

:walmartDF.show()
```

Date	Open	High	Low	Close	Volume	Adj Close
2012-01-03 00:00:00	59.970001	61.060001	59.869999	60.330002	12668800	52.619234999999996
2012-01-04 00:00:00	60.209998999999996	60.349998	59.470001	59.709998999999996	9593300	52.078475
2012-01-05 00:00:00	59.349998	59.619999	58.369999	59.419998	12768200	51.825539
2012-01-06 00:00:00	59.419998	59.450001	58.869999	59.0	8069400	51.45922
2012-01-09 00:00:00	59.029999	59.549999	58.919998	59.18	6679300	51.616215000000004
2012-01-10 00:00:00	59.43	59.709998999999996	58.98	59.040001000000004	6907300	51.494109

## 3. What does the Schema look like?

```
In [7]: walmartDF.printSchema()

root
 |-- Date: timestamp (nullable = true)
 |-- Open: double (nullable = true)
 |-- High: double (nullable = true)
 |-- Low: double (nullable = true)
 |-- Close: double (nullable = true)
 |-- Volume: integer (nullable = true)
 |-- Adj Close: double (nullable = true)
```

## 4. Print out the first 5 columns.

```
In [11]: walmartDF.show(5)
```

Date	Open	High	Low	Close	Volume	Adj Close
2012-01-03 00:00:00	59.970001	61.060001	59.869999	60.330002	12668800	52.619234999999996
2012-01-04 00:00:00	60.209998999999996	60.349998	59.470001	59.709998999999996	9593300	52.078475
2012-01-05 00:00:00	59.349998	59.619999	58.369999	59.419998	12768200	51.825539
2012-01-06 00:00:00	59.419998	59.450001	58.869999	59.0	8069400	51.45922
2012-01-09 00:00:00	59.029999	59.549999	58.919998	59.18	6679300	51.616215000000004

only showing top 5 rows

OR

```
In [99]: walmartDF.take(5)
Out[99]: [Row(Date=datetime.datetime(2012, 1, 3, 0, 0), Open=59.970001, High=61.060001, Low=59.869999, Close=60.330002, Volume=12668800, Adj Close=52.619234999999996),
Row(Date=datetime.datetime(2012, 1, 4, 0, 0), Open=60.209998999999996, High=60.349998, Low=59.470001, Close=59.709998999999999, Volume=9593300, Adj Close=52.078475),
Row(Date=datetime.datetime(2012, 1, 5, 0, 0), Open=59.349998, High=59.619999, Low=58.369999, Close=59.419998, Volume=12768200, Adj Close=51.825539),
Row(Date=datetime.datetime(2012, 1, 6, 0, 0), Open=59.419998, High=59.450001, Low=58.869999, Close=59.0, Volume=8069400, Adj Close=51.45922),
Row(Date=datetime.datetime(2012, 1, 9, 0, 0), Open=59.029999, High=59.549999, Low=58.919998, Close=59.18, Volume=6679300, Adj Close=51.616215000000004)]
```

## 5. Use describe() to learn about the DataFrame.

```
In [14]: walmartDF.describe().show()
```

	summary	Open	High	Low	Close	Volume	Adj Close
count		1258	1258	1258	1258	1258	1258
mean		72.35785375357709	72.83938807631165	71.9186009594594	72.38844998012726	8222093.481717011	67.23883848728146
stddev		6.76809024470826	6.768186808159218	6.744075756255496	6.756859163732991	4519780.8431556	6.722609449996857
min		56.389998999999996	57.060001	56.299999	56.419998	2094900	50.363689
max		90.800003	90.970001	89.25	90.470001	80898100	84.91421600000001

## 6. Bonus Question!

There are too many decimal places for mean and stddev in the describe() dataframe. Format the numbers to just show up to two decimal places. Pay careful attention to the datatypes that .describe() returns, we didn't cover how to do this exact formatting, but we covered something very similar.

```
In [93]: from pyspark.sql.functions import round, col
In [94]: wmt_desc = walmartDF.describe()
In [95]: wmt_desc = wmt_desc.select("summary", wmt_desc["Open"].cast('double'), wmt_desc["High"].cast('double'),
wmt_desc["Low"].cast('double'), wmt_desc["Close"].cast('double'),
wmt_desc["Volume"].cast('double'), wmt_desc["Adj Close"].cast('double'))
In [96]: wmt_desc.printSchema()
root
 |-- summary: string (nullable = true)
 |-- Open: double (nullable = true)
 |-- High: double (nullable = true)
 |-- Low: double (nullable = true)
 |-- Close: double (nullable = true)
 |-- Volume: double (nullable = true)
 |-- Adj Close: double (nullable = true)
In [128]: wmt_desc.select("summary", round("Open",2).alias("Open"), round("High",2).alias("High"), round("Low",2).alias("Low"),
round("Close",2).alias("Close"), round("Volume",2).alias("Volume"), round("Adj Close",2).alias("Adj Close")).show()
```

summary	Open	High	Low	Close	Volume	Adj Close
count	1258.0	1258.0	1258.0	1258.0	1258.0	1258.0
mean	72.36	72.84	71.92	72.39	8222093.48	67.24
stddev	6.77	6.77	6.74	6.76	4519780.84	6.72
min	56.39	57.06	56.3	56.42	2094900.0	50.36
max	90.8	90.97	89.25	90.47	8.08981E7	84.91

## 7. Create a new dataframe with a column called HV Ratio that is the ratio of the High Price versus volume of stock traded for a day.

```
In [131]: wmt_HVRatioDF = walmartDF.withColumn("HV Ratio", walmartDF.High/walmartDF.Volume)
```

```
In [133]: wmt_HVRatioDF.select("HV Ratio").show()
```

```
+-----+
|          HV Ratio|
+-----+
|4.819714653321546E-6|
|6.290848613094555E-6|
|4.669412994783916E-6|
|7.367338463826307E-6|
|8.915604778943901E-6|
|8.644477436914568E-6|
|9.351828421515645E-6|
|8.29141562102703E-6|
|7.712212102001476E-6|
|7.071764823529412E-6|
|1.015495466386981E-5|
|6.576354146362592...|
|5.90145296180676E-6|
|8.547679455011844E-6|
|8.420709512685392E-6|
|1.041448341728929...|
|8.316075414862431E-6|
|9.721183814992126E-6|
```

===== PROBLEM STATEMENTS  
=====

## 1. What day had the Peak High in Price?

```
In [45]: # walmartDF.orderBy(walmartDF['High'].desc()) \
#         .select(['Date']) \
#         .head(1)[0]['Date']

peakHighDF = spark.sql("SELECT Date FROM walmart WHERE High = (SELECT MAX(High) FROM walmart)")
peakHighDF.head()[0]['Date']
```

```
Out[45]: datetime.datetime(2015, 1, 13, 0, 0)
```

## 2. How many days was the Close lower than 60 dollars?

```
In [63]: close_lower60DF = spark.sql("SELECT Date FROM walmart WHERE Close < 60 ")
close_lower60DF.count()
```

```
Out[63]: 81
```

## 3. What percentage of the time was the High greater

than 80 dollars ?

```
In [69]: high_greater80DF = spark.sql("SELECT Date FROM walmart WHERE High > 80")
|(high_greater80DF.count() / walmartDF.count()) * 100
```

```
Out[69]: 9.141494435612083
```

## 4. What is the average Close for each Calendar Month?

```
In [81]: avgCloseDF = spark.sql("SELECT YEAR(Date) AS year, MONTH(Date) AS month, AVG(Close) AS avg_close\
FROM walmart GROUP BY YEAR(Date), month(Date)")
avgCloseDF.show(60)
```

year	month	avg_close
2012	10	75.30619061904761
2015	2	85.52315805263159
2014	4	77.80857085714285
2015	12	59.98681827272728
2016	7	73.54149939999999
2016	11	70.30476261904762
2012	8	73.04478265217392
2013	2	70.62315857894738
2012	4	60.149000150000006
2012	12	69.71100009999999
2014	10	76.48869486956522
2016	5	68.05285676190476
2014	12	85.1259102727273
2013	9	74.4395005
2013	10	74.97913104347826
2014	5	77.38095276190477
2016	2	66.24800044999999
2013	12	78.7752382857143
2014	1	76.53142833333334
2013	3	73.43649940000002
2014	8	74.6766623809525
2013	6	74.97800020000001
2013	7	73.40666661000001

## 5. What is the max High per year?

```
In [139]: maxHighDF = spark.sql("SELECT YEAR(Date), MAX(High) as maxHigh FROM walmart GROUP BY YEAR(Date)")
maxHighDF.show()
```

year(Date)	maxHigh
2015	90.970001
2013	81.370003
2014	88.089996
2012	77.599998
2016	75.190002

## 6. What is the mean of the Close column?

```
In [141]: meanCloseDF = spark.sql("SELECT AVG(close) as close_mean FROM walmart")
meanCloseDF.show()
```

close_mean
72.38844998012726

```
***** using HIVE
*****
```

### 1. Put the data in MySQL

#### 1.1 CREATE TABLE in MySQL:

```
CREATE TABLE walmart(date DATE, open DOUBLE, high DOUBLE, low DOUBLE, close DOUBLE, volume INT, adj_close DOUBLE);
```

```
mysql> CREATE TABLE walmart(date DATE, open DOUBLE, high DOUBLE, low DOUBLE, close DOUBLE, volume INT, adj_close DOUBLE);
-> ;
Query OK, 0 rows affected (0.02 sec)
```

#### 1.2 Loading file from LFS (VM) -----> MySQL:

```
LOAD DATA LOCAL INFILE '/home/cloudera/walmart_stock.csv' INTO TABLE walmart FIELDS TERMINATED BY ',';
```

```
mysql> LOAD DATA LOCAL INFILE '/home/cloudera/walmart_stock.csv' INTO TABLE walmart FIELDS TERMINATED BY ',';
Query OK, 1259 rows affected, 7 warnings (0.03 sec)
Records: 1259 Deleted: 0 Skipped: 0 Warnings: 6
```

### 2. SQOOP Pipeline for MySQL -----> Hive:

```
sqoop import --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table
```

```
walmart --hive-import -m 1
```

```
[clouderaquickstart ~]$ sqoop import --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table walmart --hive-import --hive-database BDSpark_pro -m 1
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
23/05/25 06:44:01 INFO SqoopRunner: Running Sqoop version: 1.4.6-cdh5.8.0
23/05/25 06:44:01 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
23/05/25 06:44:01 INFO tool.BaseSqoopTool: Using Hive-specific delimiters for output. You can override
23/05/25 06:44:01 INFO tool.BaseSqoopTool: delimiters with --fields-terminated-by, etc.
23/05/25 06:44:01 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
23/05/25 06:44:01 INFO tool.CodeGenTool: Beginning code generation
23/05/25 06:44:02 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'walmart' AS t LIMIT 1
23/05/25 06:44:02 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'walmart' AS t LIMIT 1
23/05/25 06:44:02 INFO orm.CompilationManager: HADOOP MAPRED HOME is /usr/lib/hadoop-mapreduce
Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-common-1.1.0-cdh5.8.0.jar!/hive-log4j.properties
OK
Time taken: 2.239 seconds
Loading data to table bdsark_pro.walmart
chgrp: changing ownership of 'hdfs://quickstart.cloudera:8020/user/hive/warehouse/bdsark_pro.db/walmart/part-m-00000': User does not belong to supergroup
Table bdsark_pro.walmart stats: [numFiles=1, totalSize=83365]
OK
Time taken: 1.014 seconds
```

Table in Hive:

```
Time taken: 0.004 seconds, Fetched: 1 row(s)
hive> use BDSpark_pro;
OK
Time taken: 0.024 seconds
hive> show tables;
OK
walmart
Time taken: 0.027 seconds, Fetched: 1 row(s)
hive> select * from walmart limit 5;
OK
2012-01-03      59.970001      61.060001      59.869999      60.330002      12668800      52.619235
2012-01-04      60.209998      60.349998      59.470001      59.709999      9593300 52.078475
2012-01-05      59.349998      59.619999      58.369999      59.419998      12768200      51.825539
2012-01-06      59.419998      59.450001      58.869999      59.0      8069400 51.45922
2012-01-09      59.029999      58.919998      59.18      6679300 51.616215
Time taken: 0.164 seconds, Fetched: 5 row(s)
```

### 3. Performing Queries:

## 7. What is the Pearson correlation between High and Volume?

#### 7.1-> Ext Table Creation:

```
CREATE EXTERNAL TABLE pearson_corr(pearson_corr_high_volume DOUBLE) ROW FORMAT DELIMITED FIELDS TERMINATED BY
','
```

#### 7.2-> Transferring o/p to ext table:

```
insert overwrite table pearson_corr SELECT (Avg(high * volume) - (Avg(high) * Avg(volume))) /
(stddev_pop(high) * stddev_pop(volume)) AS 'Pearsonsr' from walmart;
```

```
hive> CREATE EXTERNAL TABLE pearson_corr(pearson_corr_high_volume DOUBLE) ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
> ;
OK
Time taken: 0.530 seconds
hive> insert overwrite table pearson_corr SELECT (Avg(high * volume) - (Avg(high) * Avg(volume))) / (stddev_pop(high) * stddev_pop(volume)) AS 'Pearsonsr' from walmart;
FAILED: ParseException line 1:137 character '*' not supported here
line 1:147 character '*' not supported here
hive> insert overwrite table pearson_corr SELECT (Avg(high * volume) - (Avg(high) * Avg(volume))) / (stddev_pop(high) * stddev_pop(volume)) AS 'Pearsonsr' from walmart;
Query ID = cloudera_20230525080909_41943e2c-179e-47f3-8339-25d34a85d726
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1685904878679_0907, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1685904878679_0907/
Kill Command = /usr/lib/hadoop/bin/hadoop job *kill job_1685904878679_0907
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-05-25 08:09:45.792 Stage-1 map = 0%, reduce = 0%
```

#### 7.3-> MySQL Table Created (Client's DB):

```
CREATE TABLE pearson_corr(pearson_corr_high_volume DOUBLE);
```

```
mysql> CREATE TABLE pearson_corr(pearson_corr_high_volume DOUBLE)
-> ;
Query OK, 0 rows affected (0.04 sec)
```

#### 7.4-> Sqoop command to transfer o/p table Hive -----> MySQL:

```
sqoop export --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table
max_min_volume --export-dir /user/hive/warehouse/bdsark_pro.db/max_min_volume;
```

```
[clouderaquickstart ~]$ sqoop export --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table pearson_corr --export-dir /user/hive/warehouse/bdsark_pro
.db/pearson_corr
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
23/05/25 08:13:28 INFO SqoopRunner: Running Sqoop version: 1.4.6-cdh5.8.0
23/05/25 08:13:28 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
23/05/25 08:13:30 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
23/05/25 08:13:30 INFO tool.CodeGenTool: Beginning code generation
23/05/25 08:13:31 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'pearson_corr' AS t LIMIT 1
23/05/25 08:13:31 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'pearson_corr' AS t LIMIT 1
23/05/25 08:13:31 INFO orm.CompilationManager: HADOOP MAPRED HOME is /usr/lib/hadoop-mapreduce
```

Table in MySQL:

```
mysql> show tables;
+-----+
| Tables_in_BDSpark_pro |
+-----+
| max_min_volume         |
| pearson_corr           |
| walmart                 |
+-----+
3 rows in set (0.01 sec)

mysql> select * from pearson_corr;
+-----+
| pearson_corr_high_volume |
+-----+
| -0.338432606173703      |
+-----+
1 row in set (0.00 sec)
```

## 8. What is the max and min of the Volume column?

### 8.1-> Ext Table Creation:

```
CREATE EXTERNAL TABLE max_min_volume(maxV INT, minV INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
```

### 8.2-> Transferring o/p to ext table:

```
insert overwrite table max_min_volume select MAX(volume) as max_volume, MIN(volume) as min_volume from walmart;
```

```
hive> CREATE EXTERNAL TABLE max_min_volume(maxV INT, minV INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
>;
K
Time taken: 0.266 seconds
hive> insert overwrite table max_min_volume select MAX(volume) as max_volume, MIN(volume) as min_volume from walmart
>;
Query ID = cloudera_20230525065858_dcbfd0f0-fa99-422d-96be-ac2a6436b00f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1685004878679_0004, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1685004878679_0004/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1685004878679_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
023-05-25 06:58:27,561 Stage-1 map = 0%, reduce = 0%
Loading data to table bdspark_pro.max_min_volume
Table bdspark_pro.max_min_volume stats: [numFiles=1, numRows=1, totalSize=17, rawDataSize=16]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.9 sec HDFS Read: 91516 HDFS Write: 99 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 900 msec
OK
Time taken: 63.453 seconds
hive> show tables;
OK
max_min_volume
walmart
Time taken: 0.035 seconds, Fetched: 2 row(s)
hive> select * from max_min_volume;
OK
80898100      2094900
Time taken: 0.265 seconds, Fetched: 1 row(s)
```

### 8.3-> MySQL Table Created (Client's DB):

```
CREATE TABLE max_min_volume(max_volume INT, min_volume INT)
```

```
mysql> CREATE TABLE max_min_volume(max_volume INT, min_volume INT)
->;
Query OK, 0 rows affected (0.02 sec)
```

### 8.4-> Sqoop command to transfer o/p table Hive -----> MySQL:

```
sqoop export --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table max_min_volume --export-dir /user/hive/warehouse/bdspark_pro.db/max_min_volume;
```

```
[cloudera@quickstart ~]$ sqoop export --connect jdbc:mysql://localhost:3306/BDSpark_pro --username root --password cloudera --table max_min_volume --export-dir /user/hive/warehouse/bdspark_p
ro.db/max_min_volume;
Warning: /usr/lib/sqoop/.:/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
23/05/25 07:18:15 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.8.0
23/05/25 07:18:15 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
23/05/25 07:18:16 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
23/05/25 07:18:16 INFO tool.CodeGenTool: Beginning code generation
23/05/25 07:18:17 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'max_min_volume' AS t LIMIT 1
23/05/25 07:18:17 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'max_min_volume' AS t LIMIT 1
23/05/25 07:18:17 INFO org.apache.hadoop.mapred.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
```

Table in MySQL:

```
mysql> show tables;
+-----+
| Tables_in_BDSpark_pro |
+-----+
| max_min_volume         |
| walmart                 |
+-----+
2 rows in set (0.01 sec)

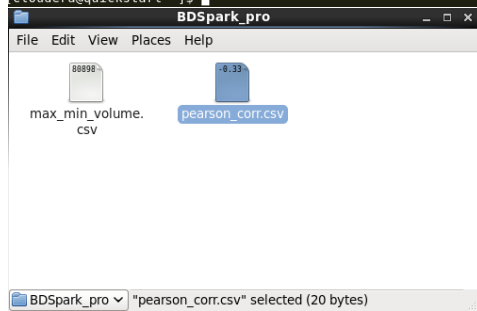
mysql> select * from max_min_volume;
+-----+
| max_volume | min_volume |
+-----+
| 80898100   | 2094900    |
+-----+
1 row in set (0.00 sec)
```

=====

#### 4. Getting o/p data files stored in HDFS -----> Client machine(here, LFS):

```
hdfs dfs -get /user/hive/warehouse/bdspark_pro.db/max_min_volume/000000_0  
/home/cloudera/BDSpark_pro/max_min_volume.csv  
hdfs dfs -get /user/hive/warehouse/bdspark_pro.db/pearson_corr/000000_0  
/home/cloudera/BDSpark_pro/pearson_corr.csv
```

```
cloudera@quickstart ~]$ hdfs dfs -get /user/hive/warehouse/bdspark_pro.db/max_min_volume/000000_0 /home/cloudera/BDSpark_pro/max_min_volume.csv  
cloudera@quickstart ~]$ hdfs dfs -get /user/hive/warehouse/bdspark_pro.db/pearson_corr/000000_0 /home/cloudera/BDSpark_pro/pearson_corr.csv  
cloudera@quickstart ~]$
```



# THANK YOU !